

Building interactive web apps

with the R package shiny

Hanne Oberman

Slides available from tinyurl.com/shiny2023 (<https://tinyurl.com/shiny2023>)

30-05-2023

What we'll discuss

1. The shiny framework
2. The user interface (UI)
3. The server
4. Advanced topics
5. Take-aways

Case study



The palmer penguins (<https://allisonhorst.github.io/palmerpenguins/>) penguins dataset (DOI [10.5281/zenodo.3960218](https://doi.org/10.5281/zenodo.3960218)).

Case study



```
library(palmerpenguins)
head(penguins)
```

```
## # A tibble: 6 × 8
##   species island    bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
##   <fct>   <fct>         <dbl>         <dbl>           <int>         <int>
## 1 Adelie  Torgersen         39.1          18.7            181          3750
## 2 Adelie  Torgersen         39.5          17.4            186          3800
## 3 Adelie  Torgersen         40.3           18            195          3250
## 4 Adelie  Torgersen          NA           NA             NA           NA
## 5 Adelie  Torgersen         36.7          19.3            193          3450
## 6 Adelie  Torgersen         39.3          20.6            190          3650
## # i 2 more variables: sex <fct>, year <int>
```

Case study

Flipper and bill length

Dimensions for Adelie, Chinstrap and Gentoo Penguins at Palmer Station LTER



The shiny framework

The basics

What is shiny?

- An R package for building shiny apps.

What is a shiny app?

- A fully interactive (web) application, which can be:
 - build as a dashboard;
 - hosted online on a webpage;
 - included in R Markdown documents.

The aim

Why use shiny?

- To create apps!
- Make your R workflows:
 - interactive (point-and-click style);
 - reproducible for non-coders;
 - look instantly professional.

The package

What does shiny offer?

- A collection of wrapper functions to write “app languages”:
 - geared toward R users who have zero experience with web development;
 - no knowledge of HTML/CSS/JavaScript required;
 - but you *can* extend it with CSS themes, HTML widgets, and JavaScript actions.
- Developed by Posit (<https://posit.co/>), so documentation and support are more or less guaranteed.

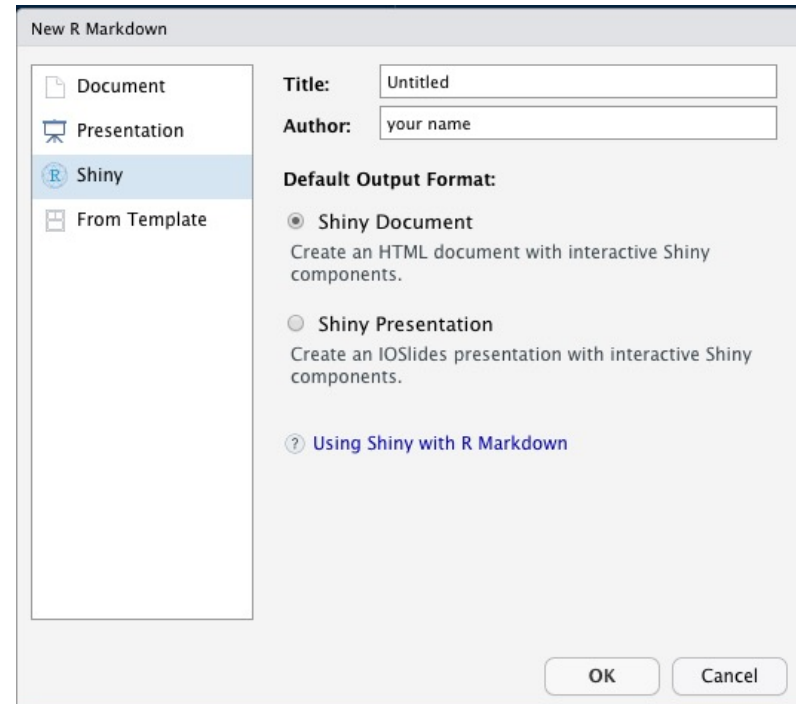
The template app

How to build a shiny app?

A. Create a file called `app.R` and add shiny components*

B. In RStudio: File → New file → R Markdown → Shiny

*file name and components are non-negotiable!




The template app


Number of bins:

20

Bandwidth adjustment:



Geyser eruption duration



The components

What does a shiny app consist of?

- A user interface (UI):
 - the visible, interactive part;
 - e.g., a web app or dashboard.
- A server:
 - the invisible, processing part;
 - e.g., your own computer or shinyapps.io (shinyapps.io).

```
library(shiny)
ui <- # some code to generate the UI
server <- # some code to generate the server
shinyApp(ui = ui, server = server)
```

Starting point

```
library(shiny)
ui <- # some code to generate the UI
server <- # some code to generate the server
shinyApp(ui = ui, server = server)
```

An empty app

```
library(shiny)
ui <- fluidPage()
server <- function(input, output) {}
shinyApp(ui = ui, server = server)
```

Our app

Recap

- Which two components does a shiny app require?
- What part of the app can be thought of as the 'front-end'?
- Do the computations need to be performed locally?

Tips

Don't rush into coding when you should be thinking.

Before building a shiny app, think about:

- What is the app aimed at?
- Who are the end users of your app? Are they tech-literate?
- In what context will the app be used? On what machines (e.g., because of screen size)?

Tips

While building a shiny app:

- KISS: Keep It Simple, Stupid;
- Use a design/UI first approach;
- Build the front-end and the back-end separately;
- If you copy something just *once*, make it a function;
- Avoid unnecessary complexity and 'feature creep'.

Tips

After building a shiny app:

- Share the app;
- Make it last.

Note. We'll get back to this later!

The UI

Starting point

```
library(shiny)
ui <- # some code to generate the UI
server <- # some code to generate the server
shinyApp(ui = ui, server = server)
```

Filling in the UI

```
ui <- fluidPage()
```

The shiny function `fluidPage()`

- generates an 'empty canvas' for shiny apps;
- en-captures all other UI elements.

Adding some text

To show text in our app, we can just add character/string objects inside `fluidPage()`:

```
ui <- fluidPage(  
  "Palmer Penguins",  
  "An interactive visualization"  
)
```

Our app

Palmer Penguins An interactive visualization

Adding formatted text

For formatted text, shiny has many functions that are wrappers around HTML tags. For example:

- `h1()`: top-level header;
- `h2()`: secondary header;
- `strong()`: bold text;
- `em()`: italicized text;
- `br()`: line break;
- `img()`: image;
- `a()`: hyperlink, etc.

Note. If you already know HTML, you don't need to use these wrapper functions!

Adding formatted text

Let's replace the UI part of our code with the following:

```
ui <- fluidPage(  
  h1("Palmer penguins"),  
  "An",  
  em("interactive"),  
  br(),  
  strong("visualization")  
)
```

Our app

Palmer penguins

An interactive
visualization

Adding a title

```
ui <- fluidPage(  
  titlePanel("Palmer penguins: An interactive visualization")  
)
```

The shiny function `titlePanel()`

- adds a visible big title-like text to the top of the page;
- sets the “official” title of the web page (i.e., displayed at the name of the tab in the browser).

Our app

Palmer penguins: An interactive visualization

Adding a layout

The simple sidebar layout:

- provides a two-column layout;
- with a smaller sidebar and a larger main panel;
- visually separates the input and output of the app.

```
ui <- fluidPage(  
  titlePanel("Palmer penguins: An interactive visualization"),  
  sidebarLayout(  
    sidebarPanel("[inputs]"),  
    mainPanel("[outputs]")  
  )  
)
```


Our app

Palmer penguins: An interactive visualization

[inputs]

[outputs]

Adding an input element

Inputs allow users to interact with a shiny app.


We've seen two types already:

- `selectInput()` creates a drop-down menu;
- `sliderInput()` creates a numeric scale.

Number of bins:

20

Bandwidth adjustment:



Adding an input element

Can you guess what kind of element these input functions will create?

- `textInput();`
- `dateInput();`
- `checkboxInput().`

Adding input elements

Button

Action

`actionButton()`

Single checkbox

☒ Choice A

`checkboxInput()`

Checkbox group

☒ Choice 1
☐ Choice 2
☐ Choice 3


`checkboxGroupInput()`

Date input

2014-01-01

`dateInput()`

Colour input



`colourpicker::colourInput()`

Date range

2014-01-24 to 2014-01-24

`dateRangeInput()`

File input

Choose File No file chosen

`fileInput()`

Numeric input

1

`numericInput()`

Password Input

`passwordInput()`

Radio buttons

☒ Choice 1
☐ Choice 2
☐ Choice 3

`radioButtons()`

Select box

Choice 1

`selectInput()`

Sliders



`sliderInput()`

Text input

Enter text...

`textInput()`

Text area

Multiple lines
of text

`textAreaInput()`

Adding an input element

Which function would suit an input element for the variable `island`?

Location
Biscoe ☒
Dream ☐
Torgersen ☐

```
radioButtons(  
  inputId = "location",  
  label = "Location",  
  choices = c("Biscoe", "Dream", "Torgersen")  
)
```

Adding an input element

All input functions have the same first two arguments:

- `inputId`, the name by which shiny will refer to this input when you want to retrieve its current value;
- `label`, which specifies the text displayed right above the input element.

These argument names are typically dropped from the `...Input()` function call:

```
radioButtons("location", "Location", choices = c("Biscoe", "Dream", "Torgersen"))
```

Note. Every input in your app *must* have a unique `inputId`; the app will not work properly otherwise! So keep your `inputIds` simple and sensible.

Adding an input element

The resulting UI code looks like:

```
ui <- fluidPage(  
  titlePanel("Palmer penguins: An interactive visualization"),  
  sidebarLayout(  
    sidebarPanel(  
      radioButtons("location", "Location",  
                   choices = c("Biscoe", "Dream", "Torgersen"))  
    ),  
    mainPanel("[outputs]")  
  )  
)
```

Our app

Palmer penguins: An interactive visualization

Location

☒ Biscoe


☐ Dream

☐ Torgersen

[outputs]

Adding another input element

Let's create an input element for the variable `body_mass_g` as well. Which input function(s) should we use?



```
sliderInput("weight", "Weight",  
           min = 2, max = 7, value = c(3, 5), post = "kg")
```

Adding another input element

The full UI code is now:

```
ui <- fluidPage(  
  titlePanel("Palmer penguins: An interactive visualization"),  
  sidebarLayout(  
    sidebarPanel(  
      radioButtons("location", "Location",  
                   choices = c("Biscoe", "Dream", "Torgersen")),  
      sliderInput("weight", "Weight",  
                  min = 2, max = 7, value = c(3, 5), post = "kg")  
    ),  
    mainPanel("[outputs]")  
  )  
)
```

Our app

Palmer penguins: An interactive visualization

Location

☒ Biscoe


☐ Dream

☐ Torgersen

Weight

2kg 3kg 5kg 7kg

2 3 4 5 6 7

An interactive slider for penguin weight. It features a horizontal track with a blue line indicating the current range from 3 to 5 kg. Two circular handles are positioned at the 3 kg and 5 kg marks. The track is labeled with '2kg', '3kg', '5kg', and '7kg' at the top, and numerical values '2', '3', '4', '5', '6', and '7' at the bottom.

[outputs]

Adding an output element

Outputs are *shown* in the UI, but *created* on the server side.

That's why we add placeholders for the outputs in the UI.

Placeholders:

- Determine where an output will be;
- Give outputs a unique ID to link it to the server;
- Won't actually show anything, yet.

```
mainPanel(  
  "[plot placeholder]", plotOutput("scatterplot")  
)
```

Our app

Palmer penguins: An interactive visualization



[plot placeholder]

Adding another output element

The placeholder doesn't show anything, because we haven't created any figure yet on the server side.

But first, let's add another output element:

```
mainPanel(  
  "[plot placeholder]", plotOutput("scatterplot"),  
  br(),  
  br(),  
  "[table placeholder]", tableOutput("descriptives")  
)
```

Note. We added a few line breaks `br()` between the two outputs, so that they aren't crammed on top of each other.

The complete UI

```
ui <- fluidPage(  
  titlePanel("Palmer penguins: An interactive visualization"),  
  sidebarLayout(  
    sidebarPanel(  
      radioButtons("location", "Location",  
                   choices = c("Biscoe", "Dream", "Torgersen")),  
      sliderInput("weight", "Weight",  
                  min = 2, max = 7, value = c(3, 5), post = "kg")  
    ),  
    mainPanel(  
      "[plot placeholder]", plotOutput("scatterplot"),  
      br(),  
      br(),  
      "[table placeholder]", tableOutput("descriptives")  
    )  
  )  
)
```

Our app

Palmer penguins: An interactive visualization

Location

☒ Biscoe

☐ Dream

☐ Torgersen

Weight

2kg3kg5kg7kg

234567

[plot placeholder]

Recap

- How can formatted text be shown in a shiny app?
- What is mandatory *and* unique in input element functions?
- How do you define where an output element will be shown?

Tips

When building the front-end of your app:

- Work on the general appearance first, anything that does not rely on computation (e.g., tabs, inputs, outputs);
- Use mock data and/or text (build an 'ipsum-app');
- Make the app self-evident; the main usage of the app should not require reading any manual.

The server

Starting point

```
library(shiny)
ui <- # some code to generate the UI
server <- # some code to generate the server
shinyApp(ui = ui, server = server)
```

The server function

```
server <- function(input, output) {}
```

The server function:

- requires* input and output IDs from the UI;
- builds output objects via `render...()` functions;
- saves the generated output into an output list.


*exceptions apply!

Building static output

Let's use the exception to the rule to develop our server step-by-step.

```
ggplot(penguins,  
  aes(x = flipper_length_mm, y = bill_length_mm, color = species)) +  
  geom_point() +  
  geom_smooth(method = "lm", se = FALSE) +  
  scale_color_palmer() +  
  labs(title = "Flipper and bill length by species (static)")
```

Building static output


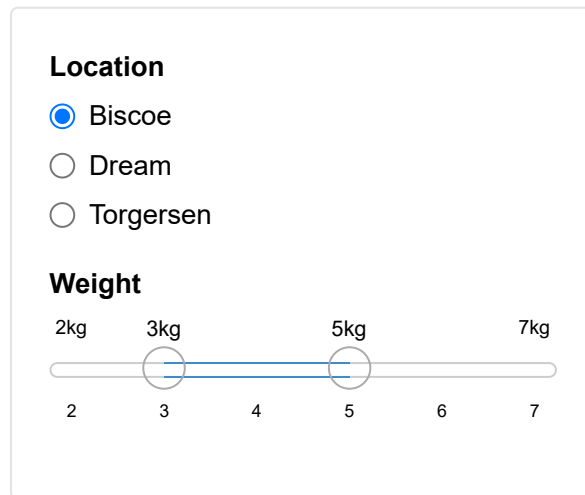


Building static output

```
server <- function(input, output) {  
  output$scatterplot <- renderPlot({  
    ggplot(penguins,  
      aes(x = flipper_length_mm, y = bill_length_mm, color = species)) +  
      geom_point() +  
      geom_smooth(method = "lm", se = FALSE) +  
      scale_color_palmer() +  
      labs(title = "Flipper and bill length by species (static)")  
  })  
}
```

Our app

Palmer penguins: An interactive visualization



Building interactive output



To make the figure interactive, we have to link the server to the UI inputs.

```
server <- function(input, output) {  
  output$scatterplot <- renderPlot({  
    ggplot(penguins,  
      aes(x = flipper_length_mm, y = bill_length_mm, color = species)) +  
    geom_point() +  
    geom_smooth(method = "lm", se = FALSE) +  
    scale_color_palmer() +  
    labs(title = paste("Flipper and bill length at", input$location))  
  })  
}
```

Whenever the selected location changes, the plot title will update.

Our app

Palmer penguins: An interactive visualization




Building an interactive visualization

To incorporate interactivity in the visualization, we have to filter the data based on the input values.

```
server <- function(input, output) {  
  output$scatterplot <- renderPlot({  
    filtered <- penguins %>%  
      filter(island == input$location,  
             body_mass_g >= input$weight[1] * 1000,  
             body_mass_g <= input$weight[2] * 1000  
            )  
    filtered %>%  
      ggplot(aes(x = flipper_length_mm, y = bill_length_mm, color = species)) +  
        geom_point() +  
        geom_smooth(method = "lm", se = FALSE) +  
        scale_color_palmer() +  
        labs(title = "Flipper and bill length")  
  })  
}
```

Our app

Palmer penguins: An interactive visualization



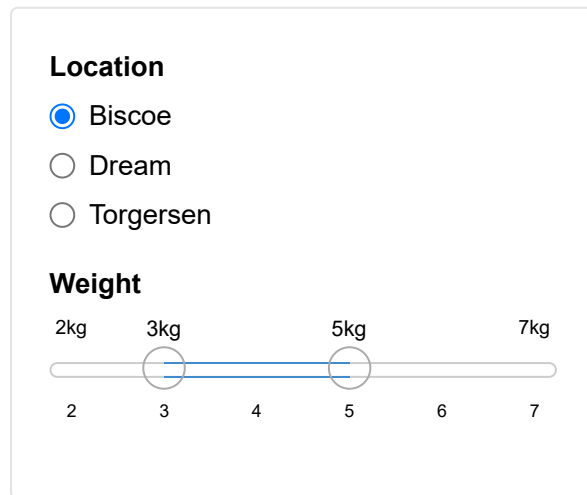
Building an interactive summary table

To complete our app we need to build an interactive output for the table placeholder and add it to the server:

```
output$descriptives <- renderTable({  
  filtered <- penguins %>%  
    filter(island == input$location,  
           body_mass_g >= input$weight[1] * 1000,  
           body_mass_g <= input$weight[2] * 1000  
    )  
  filtered %>%  
    group_by(species) %>%  
    summarise(mean_bill = mean(bill_length_mm, na.rm = TRUE),  
              mean_flip = mean(flipper_length_mm, na.rm = TRUE))  
})
```

Our app

Palmer penguins: An interactive visualization



Intermezzo

Reactivity

Short break from our app to talk about a crucial concept in shiny: reactivity.

Reactivity enables your outputs to react to changes in inputs.

On the most basic level, it means that when the value of a variable x changes, anything that relies on x (i.e. has x in it) gets re-evaluated.

Consider the following code

```
x <- 5  
y <- x + 1  
x <- 10
```

What is the value of y ?

Reactivity

What is the value of y?

```
x <- 5  
y <- x + 1  
x <- 10
```

In ordinary programming, the value of y is still 6.

In reactive programming, however, x and y are *reactive expressions*. Now, the value of y updates reactively, and becomes 11.

Reactivity is the foundation for the responsiveness of shiny apps.

Reactivity

In our server, we implicitly use reactivity when we filter the data for our outputs:

```
filtered <- penguins %>%  
  filter(island == input$location,  
         body_mass_g >= input$weight[1] * 1000,  
         body_mass_g <= input$weight[2] * 1000  
  )  
)
```

Whenever one of the inputs changes, our outputs change with it. But, this part of code is duplicated, because we didn't use a reactive variable.

Reactivity

We can avoid code duplication by:

- defining a reactive variable that will hold the filtered dataset;
- using that variable in the `render...()` functions.

```
filtered <- reactive({  
  penguins %>%  
    filter(island == input$location,  
           body_mass_g >= input$weight[1] * 1000,  
           body_mass_g <= input$weight[2] * 1000  
    )  
})
```


Reactivity

What is going on behind the scenes?

- The location input changes →
- shiny 'looks' at the reactive(s) that depend on location →
- `filtered()` is re-evaluated →
- shiny 'looks' at the reactive(s) that depend on `filtered()` →
- The two `render...()` functions are re-executed →
- The plot and the table output are updated.

This can be visualized in a dependency tree, to show what value depends on what other value.

Reactivity



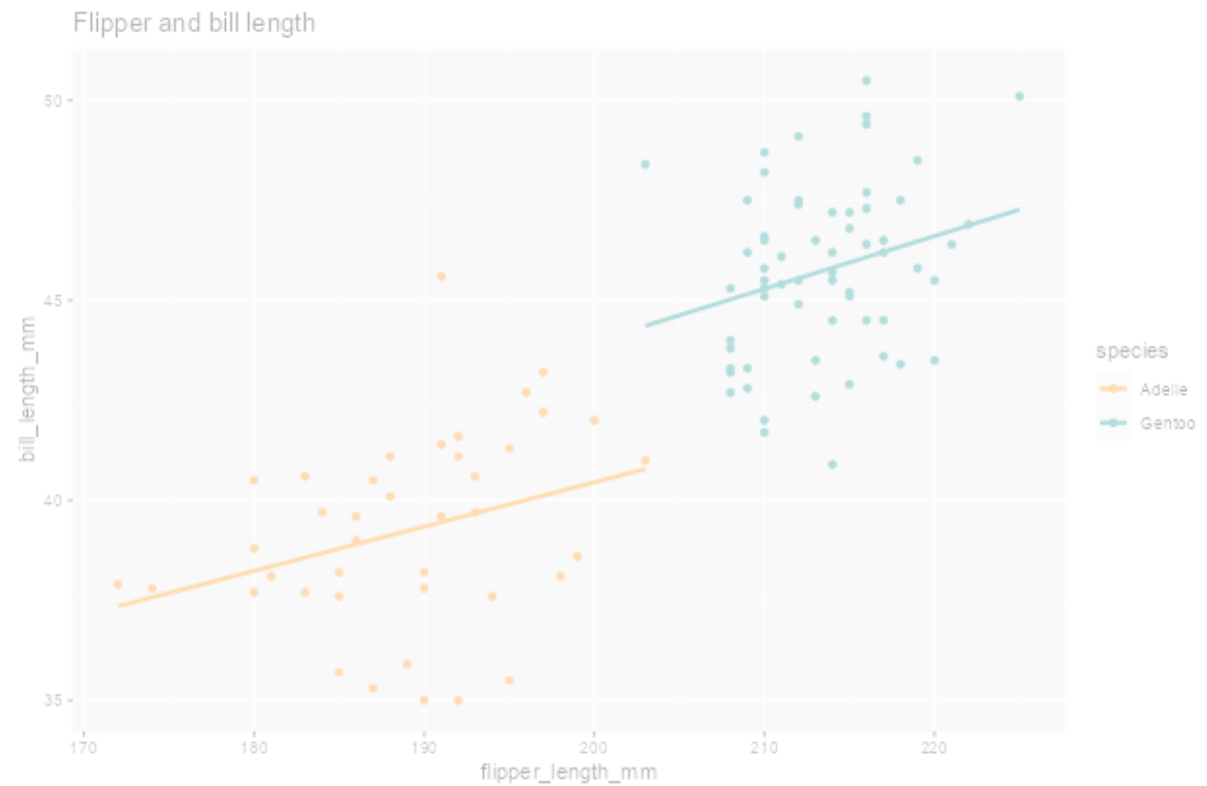
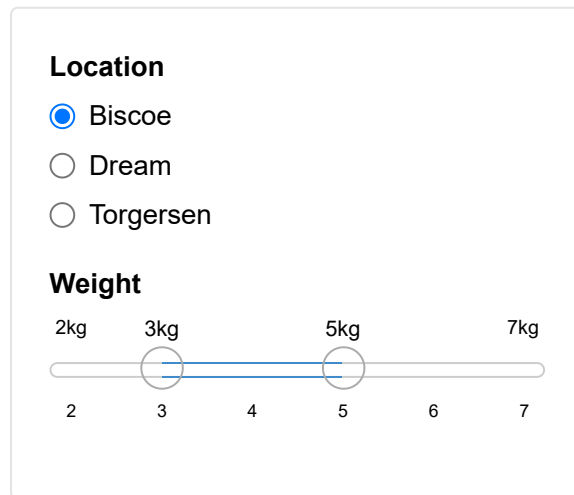
The server (continued)

The final app

```
server <- function(input, output) {  
  filtered <- reactive({  
    penguins %>%  
      filter(island == input$location,  
             body_mass_g >= input$weight[1] * 1000,  
             body_mass_g <= input$weight[2] * 1000  
            )  
  })  
  
  output$scatterplot <- renderPlot({  
    filtered() %>%  
      ggplot(aes(x = flipper_length_mm, y = bill_length_mm, color = species)) +  
        geom_point() +  
        geom_smooth(method = "lm", se = FALSE) +  
        scale_color_palmer() +  
        labs(title = "Flipper and bill length")  
  })  
  
  output$descriptives <- renderTable({  
    filtered() %>%  
      group_by(species) %>%  
      summarise(mean_bill = mean(bill_length_mm, na.rm = TRUE),  
                mean_flip = mean(flipper_length_mm, na.rm = TRUE))  
  })  
}
```

Our app

Palmer penguins: An interactive visualization



Recap

- What is the common naming element in the set of shiny output functions?
- How can you make the server side more efficient and easier to debug?
- How are the UI and server sides linked?

Tips

When building the back-end of your app:

- Use sensible non-reactive defaults while developing (e.g., `data <- mtcars` instead of `data <- reactive(...)`);
- Think about what could be 'hard coded' in the final app too, because of the reactivity vs. speed trade-off;
- Extract the complex (but non-reactive) processing functions and put them in separate files;
- Add user feedback to make server-side requirements explicit (e.g., input validation, pop-up messages, loading icons).

Advanced topics

Design

- Use more complex layouts, such as tabs or dashboards;
- Make the output elements 'clickable' with `plotly` and `datatable`;
- Change input element options from the server side with `update...()` functions.

Robustness

- Run the app in the viewer panel, a separate window, and your browser;
- Monkey test it (i.e., click EVERYTHING);
- Provide the wrong inputs (e.g., a corrupt data file, a file with the 'wrong' extension, an 'impossible' numeric input, etc.);
- Modularize your app;
- Use the `golem` framework for production-grade shiny apps (but decide upfront!).

Deployment

Deploy your app on [shinyapps.io](https://www.shinyapps.io/) (<https://www.shinyapps.io/>):

- You'll have a link to use/share the app online;
- Non-R-users will be able to interact with your app;
- You can tweak your app to cache certain outputs, or have several users in one session (like Google Drive documents);
- But, with a free account, your app will be public;
- And if your app is too popular, you will eventually need to pay server costs.

Note. You could also host your app on your own website. Or don't deploy it at all (e.g., for privacy reasons).

Take-aways

Summary

- shiny allows you to build interactive (web) apps from R;
- shiny apps consist of two parts, the user interface (UI) and the server:
 - In the UI, you design what is shown to the user,
 - In the server, you do all the modeling and building of the outputs,
 - You link the UI and the server to make the app interactive,
 - To optimize these interactions, you can use reactive expressions;
- This is only the tip of the iceberg, there are many more things you can do with shiny.

Inspiration


Check out these amazing resources:

- RStudio's introduction to shiny webinar (<https://www.rstudio.com/resources/webinars/introduction-to-shiny/>);
- Hadley Wickham's book Mastering Shiny (<https://mastering-shiny.org/>);
- The official cheatsheet (<https://rstudio.com/resources/cheatsheets/>);
- The more advanced Engineering Shiny (<https://engineering-shiny.org/>);
- This webinar on Modularizing Shiny (<https://www.youtube.com/watch?v=yILLVo2VL50>).

And look for examples here:

- The Shiny Gallery (<https://shiny.rstudio.com/gallery/>);
- The annual shiny contest (<https://www.rstudio.com/blog/winners-of-the-2nd-shiny-contest/>).

Check out my app!



hanneoberman.shinyapps.io/shinymice-demo/
(<https://hanneoberman.shinyapps.io/shinymice-demo/>)

Thank you!